

Robust estimation

* Naive approach:

- fit model to all points
- compute distance of each point from model
- Discard points with largest distance
- fit model to remaining points

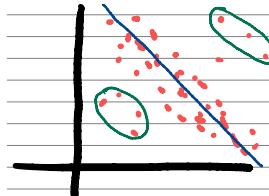
* Initial model is inaccurate and so we drift in the wrong direction.

* Approaches:

- M-estimators → modify the loss function to reduce the impact of outliers. Outliers are not discarded, their weights are down.
- RANSAC → Use a subset of data points and checks which points fit well. Repeat to find the best model with the highest number of inliers

○ → outliers, they are removed

This model can be inaccurate because we remove outliers. These outliers can be valid data points and their elimination could lead to a wrong direction



M-estimators

- Mean square Error (MSE) fitting:

$$E(\theta) = \sum d^2(x_i; \theta)$$

e.g. $d^2 = (\ell^T x_i)^2$
for line fitting

- Robust estimation:

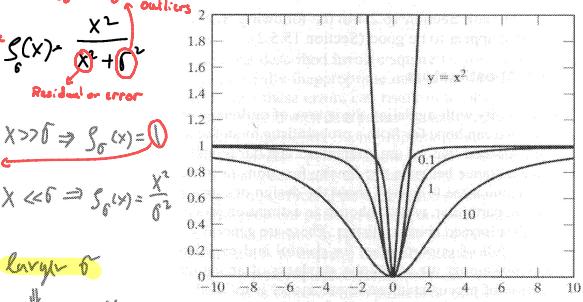
$$E(\theta) = \sum \rho(d(x_i; \theta))$$

Loss function that reduce the impact of outliers

MSE is a special case where $\rho(x) = x^2$

Geman-McClure estimator: Type of loss function to reduce the influence of outliers

Parameter that controls how aggressively the estimator treats outliers



larger δ reduce the penalty for larger residuals

Geman-McClure estimator

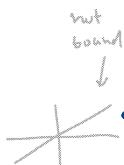
$$E(\theta) = \sum_i S_\rho(d(x_i; \theta))$$

we try to minimize for robust estimation using its gradient

gradient $\rightarrow \nabla E(\theta) = \sum \left(\frac{\partial}{\partial d} S_\rho(d) \right) \frac{\partial d(\theta)}{\partial \theta}$

For $S_\rho(d) = d^2 \Rightarrow \frac{\partial S_\rho}{\partial d} = 2d$
(MSE case)

For $S_\rho = \frac{d^2}{d+r^2} \Rightarrow \frac{\partial S_\rho}{\partial d} = \frac{2d}{(d+r^2)^2}$
Geman-McClure function



wkt bound
↓
Gradients grow linearly with d, meaning large residuals (outliers) have a big influence on the overall model.

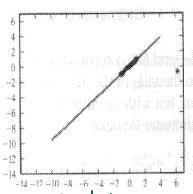


Gradient becomes smaller, outliers are down-weighted and have less impact compare to least square.

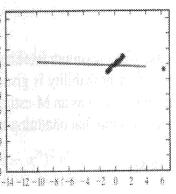
* $\frac{\partial}{\partial d} P_S(d) \Rightarrow$ Represents how the loss function P_S changes with respect to the residual d

* $\frac{\partial}{\partial \theta} d(\theta) \Rightarrow$ Represents the residual changes with respect to the model parameters

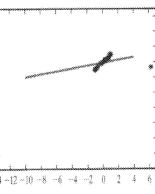
Robust estimation example



Correct sigma
Right sigma
Capture the trend of the inliers while ignoring the outlier



Small sigma
The model is sensitive to small residuals (inliers are treated as potential outliers)



Large sigma
All point are treated as inliers

Selecting bandwidth parameter

- Large $\hat{r} \Rightarrow$ include more points
- small $\hat{r} \Rightarrow$ include fewer points
- Variable estimation:
start with large \hat{r} and decrease as converging

$$\hat{r}^{(n)} = 1.5 \text{ median } \{ d(x_i; \theta^{(n-1)}) \}$$

\uparrow
Estimate at step n \uparrow
Parameters at step n-1

M-estimator summary

- 1) Draw a large set of points uniformly at random
- 2) Select initial value of θ
- 3) Fit model $\rightarrow \theta^{(i)}$
- 4) Compute $\sigma^{(i)}$ using median distance of points
- 5) Continue while objective is decreasing

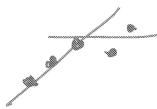
* To overcome incorrect initial guess of θ
 repeat several times and select best solution
 (smaller objective)

- ① Select a subset
- ② Select a σ (large $\sigma \rightarrow$ more points)
- ③ Fit the model to get θ (based on minimizing M-estimator loss function)
- ④ Compute the σ new using the median distance of points
- ⑤ Continue until the objective is decreasing

RANSAC

Random Sample Consensus (RANSAC):

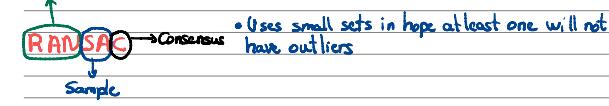
- perform multiple experiments
- choose best results
- use small sets in hope that at least one set will not have outliers



Parameters:

- n = # points drawn at each evaluation \rightarrow Define the size of subset
- d = min # points needed to estimate model \rightarrow Estimate a line \rightarrow At least 2 points
- K = # trials
- t = distance threshold to identify inliers

- ① Select subset ③ Choose the best Explain the majority of outliers
② Train this subset Does not have outliers (hope)



RANSAC Algorithm

- * Repeat K times:
 - Draw n points uniformly at random (with replacement)
 - fit a model to points
 - Find inliers in entire set (distance $< t$)
 - Recompute model (if at least d inliers)
 - update parameters (K, t)
- * Choose best solution:
 - Largest consensus set (- or smallest error)

- ① Select subset
- ② Fit a model to the points
- ③ Find inliers (distance $< t$)
- ④ Recompute the model (if d inliers are found) \rightarrow Improve the accuracy
- ⑤ Update K and t

Best model
Less error (σ_r)
Best consensus set

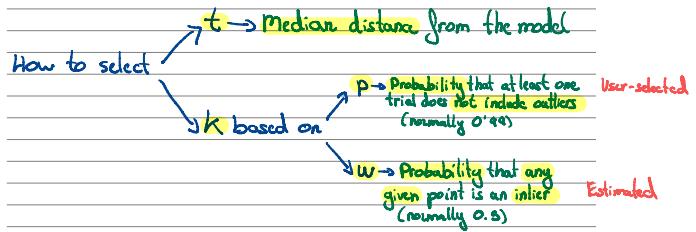
Estimating RANSAC parameters

* To estimate t use median distance from model.

* To estimate k use:

P : with probability of p at least one experiment does not have outliers (e.g. $P \approx 0.99$) ^{user selected}

w : probability that a point is an inlier (initially $w=0.5$) ^{estimated}



Estimating RANSAC parameters

Probability that all k experiments failed:

$$(1-p) = (1-w^k)^k$$

Probability that at least one trial succeeds (no outliers)

$$\log(1-p) = k \log(1-w^k) \leftarrow \text{Taking log of both sides...}$$

$$k = \frac{\log(1-p)}{\log(1-w^k)}$$

large $p \rightarrow$ large k A higher confidence level requires more iterations
small $w \rightarrow$ large k If w is small, means there are many outliers

$$w \leftarrow \frac{\# \text{ inliers}}{\# \text{ points}}$$

Update w in every iteration but set upper bound for k

Segmentation

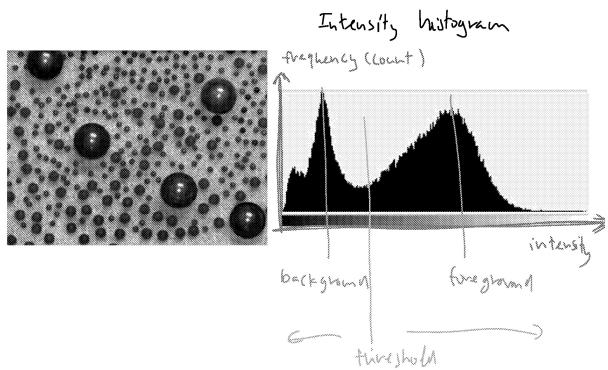
* Split image into sub-parts:

- 1) Based on color
- 2) Based on spatial location
- 3) Based on features
- 4) Based on semantics

Problem statement

- - Separate object(s) from background
 - find contours of objects
 - Semantic image segmentation: label each pixel in the image with class label

Color segmentation



Spatial context

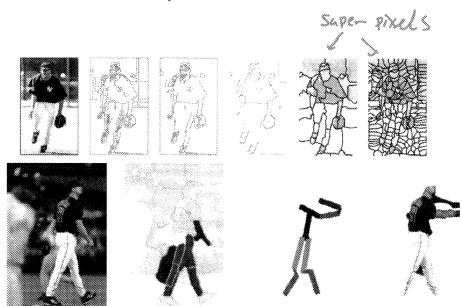


FIGURE 9.14: Superpixels can expose structure in images that other representations conceal (e.g., human body segments tend to appear as long, thin segments). **(a)** On the top row, an image together with three different edge maps, and superpixels computed at two “scales”. Note that the coarser superpixels tend to expose limb segments. **(b)** On the bottom row, another image, its superpixels, and two versions of the body layout inferred from the superpixel representation.

Feature based segmentation

* Define feature vector at each pixel x:

$$F(x) = \begin{bmatrix} x \\ I(x) \\ L(x) \end{bmatrix} \begin{array}{l} \leftarrow \text{location} \\ \leftarrow \text{intensity} \\ \leftarrow \text{local characteristics} \\ \quad \quad \quad \text{(e.g. texture)} \end{array}$$

* Apply clustering

Clustering

* Algorithms:

- k-means
 - Mixtures of Gaussians
 - Mean shift
 - Expectation maximization
 - Graph cuts
 - spectral clustering